

Prelim Notes

Dylan

Eric

Robert

Shubham

1. Natural Language Processing and Speech

1.1. Chapter 1-4 (July 18)

Discussions

- Are n-grams defined over word forms or word roots ?
Ans - application specific.
- Good-Turing smoothing : we derived Eq 4.27 using Eq 4.26 and briefly discussed the approximation used for larger N (Simple-Good turing)
- Brief discussion of interpolation and katz-backoff.
- Discussed back-off in Kneser Ney. An unanswered question was regarding implementation of back-off from n-grams to (n-1)-grams for $n > 1$ (do we use context or back-off to Kneser-Ney probabilities ?)

Topics for Review

- Kneser-Ney
- Perplexity
- Good Turing (formula)

1.2. Chapter 5-9 (July 22)

Discussions

- Maximum entropy Markov model (how this differs from hidden Markov model)
- Hidden Markov model speech recognition (how the model works)

Topics for Review

- Part of speech tagging (hidden Markov models and rule based)
- Maximum entropy and maximum entropy markov models
- features for speech (pitch, amplitude, spectrograms, formants)
- Cepstrum
- Speech recognition with hidden Markov model

1.3. Chapter 10, 12-13 (July 25)

Discussions

- Viterbi Approximation
- Only simple language models
- multi-pass – do coarse solution proposal, exact comparison of proposals
- word lattices – not 100% on how to build it, but its probably just a modified viterbi
- confidence
- A* decoding – fast match; $g(\text{cost}) + \text{partial path}$ is confusing
- triphones – bigram to trigram model; state space expansion; decision tree
- MMIE – wtf
- adaptation – gender; Maximum Likelihood Linear Regression (MLLR)
- Dependency vs CFG
- Bottom-up vs top-down
- parsing hard b/c ambiguity
- CKY

Topics for Review

- Stack/ A* decoding
- Multi-Pass
- Tri-Phones
- Adaptation
- Know Context Free Grammars
- Chomsky Normal Form

2. Artificial Intelligence: A Modern Approach

2.1. Chapter 1-5 (July 22)

Discussions

- variants of A^* (limited memory, weighted, $LRTA^*$)
- And-or search trees (their different uses)
- alpha-beta pruning (work through an example)

Topics for Review

- BFS, DFS, A^* , limited-memory A^*
- Iterative deepening DFS
- Heuristics for A^*
- Simulated annealing, evolutionary algorithms
- And-or search trees
- $LRTA^*$
- Minimax
- Alpha-beta pruning
- Expectimax

2.2. Chapter 6-9 (July 25)

Discussions

- K-consistency – generalization of individual consistency
- Reduction to binary constraints – introduce variables, dual graphs
- path consistency – remove pairs or singletons
- backjumping – completeness; backjumping undoes any variables that were set after the node you back-jump to
- Forward Checking \subset Arc Consistency
- 6.5 – favorite section for questions
 - Tree Structured CSP – enforce directed arc consistency \rightarrow no backtracking needed
 - Cut sets – iterate over all possible assignments to the cut set; cut set leaves remaining graph as tree
 - Tree decomposition – junction tree for CSPs! – graph with subsets of variables, tree structured solver enforces variable consistency.
- Basic logic modus ponens/and resolution

- models \rightarrow possible worlds, simplest inference = model-checking – $\alpha \implies \beta$ iff $M(\alpha) \subseteq M(\beta)$.
- Forward Chaining – Horn Clauses + Definite Clauses; important idea is the lhs of implications are positive conjunctions. That property saves backtracking over incorrect assignments.
- DPLL – DFS with Heuristics; what do we need to know
- Satisfiability Problems are usually easy
- frame problem, fluents,
- belief states – represent 'belief' as things that could be true
- SATPlan – be super verbose or problems arise
- object/relations/functions
- quantifiers \exists, \forall

Topics for Review

- CSPs – know them
- consistency: node, arc, path, strong-k
- heuristics: minimum remaining value, degree, least constrained
- backjumping/conflict directed
- local search
- Tree Structured- CSPs
- knowledge base
- syntax
- possible worlds/models
- entailment
- horn clauses/forward chaining
- (logical) state estimation